

Importing Libraries

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import plotly.express as px
from tqdm import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
```

Downloading the training and testing dataset

```
!wget https://archive.ics.uci.edu/ml/machine-learning-databases/statlog/satimage/sat.trn
!wget https://archive.ics.uci.edu/ml/machine-learning-databases/statlog/satimage/sat.tst
```

```
--2021-04-03 08:48:59-- https://archive.ics.uci.edu/ml/machine-learning-databases/statlog/satimage/sat.trn
Resolving archive.ics.uci.edu (archive.ics.uci.edu)... 128.195.10.252
Connecting to archive.ics.uci.edu (archive.ics.uci.edu)|128.195.10.252|:443... connected
HTTP request sent, awaiting response... 200 OK
Length: 525830 (514K) [application/x-httpd-php]
Saving to: 'sat.trn.2'
```

```
sat.trn.2          100%[=====>] 513.51K  --.-KB/s    in 0.1s
```

```
2021-04-03 08:48:59 (3.61 MB/s) - 'sat.trn.2' saved [525830/525830]
```

```
--2021-04-03 08:48:59-- https://archive.ics.uci.edu/ml/machine-learning-databases/statlog/satimage/sat.tst
Resolving archive.ics.uci.edu (archive.ics.uci.edu)... 128.195.10.252
Connecting to archive.ics.uci.edu (archive.ics.uci.edu)|128.195.10.252|:443... connected
HTTP request sent, awaiting response... 200 OK
Length: 236745 (231K) [application/x-httpd-php]
Saving to: 'sat.tst.2'
```

```
sat.tst.2          100%[=====>] 231.20K  --.-KB/s    in 0.1s
```

```
2021-04-03 08:48:59 (1.95 MB/s) - 'sat.tst.2' saved [236745/236745]
```



Reading and Storing training and testing dataframe

```
train = pd.DataFrame(np.genfromtxt('sat.trn'))
train.rename(columns={train.columns[-1]: "Class" }, inplace=True)

print(train.head(20))
```

```
      0      1      2      3      4  ...  32      33      34      35  Class
0  92.0  115.0  120.0  94.0  84.0  ...  84.0  107.0  113.0  87.0    3.0
```

```

1  84.0  102.0  106.0  79.0  84.0  ...  84.0  99.0  104.0  79.0  3.0
2  84.0  102.0  102.0  83.0  80.0  ...  84.0  99.0  104.0  79.0  3.0
3  80.0  102.0  102.0  79.0  84.0  ...  84.0  103.0  104.0  79.0  3.0
4  84.0  94.0  102.0  79.0  80.0  ...  79.0  107.0  109.0  87.0  3.0
5  80.0  94.0  98.0  76.0  80.0  ...  79.0  107.0  109.0  87.0  3.0
6  76.0  102.0  106.0  83.0  76.0  ...  79.0  103.0  104.0  79.0  3.0
7  76.0  102.0  106.0  87.0  80.0  ...  79.0  95.0  100.0  79.0  3.0
8  76.0  89.0  98.0  76.0  76.0  ...  79.0  87.0  93.0  71.0  4.0
9  76.0  94.0  98.0  76.0  76.0  ...  79.0  87.0  93.0  67.0  4.0
10 76.0  98.0  102.0  72.0  76.0  ...  75.0  87.0  96.0  71.0  4.0
11 72.0  94.0  90.0  72.0  72.0  ...  71.0  87.0  89.0  67.0  4.0
12 72.0  89.0  94.0  76.0  72.0  ...  71.0  79.0  81.0  62.0  4.0
13 76.0  94.0  98.0  76.0  72.0  ...  67.0  75.0  85.0  62.0  4.0
14 68.0  85.0  86.0  68.0  68.0  ...  71.0  75.0  81.0  62.0  4.0
15 68.0  89.0  86.0  72.0  68.0  ...  67.0  75.0  85.0  71.0  4.0
16 68.0  85.0  90.0  76.0  68.0  ...  67.0  75.0  96.0  79.0  4.0
17 68.0  94.0  94.0  79.0  76.0  ...  75.0  83.0  96.0  83.0  4.0
18 80.0  94.0  102.0  83.0  80.0  ...  84.0  99.0  109.0  87.0  3.0
19 88.0  106.0  115.0  87.0  88.0  ...  88.0  107.0  109.0  83.0  3.0

```

```
[20 rows x 37 columns]
```

```

test = pd.DataFrame(np.genfromtxt('sat.tst'))
test.rename(columns={test.columns[-1]: "Class" }, inplace=True)

```

```
print(test.sample(20))
```

```

      0      1      2      3      4  ...  32      33      34      35  Class
672  68.0  77.0  90.0  72.0  68.0  ...  67.0  68.0  89.0  79.0    5.0
715  43.0  29.0  113.0  114.0  43.0  ...  40.0  32.0  100.0  107.0    2.0
838  64.0  73.0  74.0  57.0  64.0  ...  67.0  72.0  74.0  58.0    7.0
925  89.0  106.0  110.0  83.0  82.0  ...  76.0  87.0  91.0  67.0    4.0
1158 66.0  109.0  122.0  100.0  66.0  ...  63.0  111.0  124.0  101.0    1.0
1916 63.0  66.0  90.0  79.0  63.0  ...  88.0  111.0  120.0  94.0    1.0
1651 72.0  85.0  98.0  79.0  64.0  ...  51.0  51.0  81.0  79.0    5.0
1971 60.0  85.0  94.0  79.0  60.0  ...  63.0  79.0  100.0  87.0    1.0
1605 67.0  95.0  105.0  86.0  67.0  ...  76.0  89.0  115.0  94.0    5.0
1897 52.0  56.0  80.0  74.0  59.0  ...  57.0  67.0  85.0  76.0    5.0
28   84.0  98.0  102.0  79.0  80.0  ...  71.0  75.0  85.0  67.0    7.0
1584 66.0  113.0  127.0  100.0  66.0  ...  63.0  106.0  114.0  90.0    1.0
1848 63.0  91.0  100.0  83.0  67.0  ...  63.0  91.0  101.0  86.0    1.0
1967 76.0  106.0  115.0  94.0  76.0  ...  79.0  103.0  123.0  100.0    1.0
200  88.0  103.0  113.0  85.0  88.0  ...  86.0  100.0  108.0  81.0    3.0
467  97.0  115.0  119.0  94.0  93.0  ...  76.0  91.0  96.0  74.0    4.0
1777 44.0  31.0  114.0  140.0  44.0  ...  46.0  39.0  108.0  114.0    2.0
559  63.0  64.0  85.0  67.0  67.0  ...  67.0  66.0  72.0  53.0    7.0
1387 76.0  87.0  91.0  63.0  80.0  ...  66.0  79.0  76.0  59.0    4.0
1631 88.0  106.0  106.0  87.0  88.0  ...  93.0  107.0  109.0  87.0    3.0

```

```
[20 rows x 37 columns]
```

```

x_test = test[test.columns[:-1]].to_numpy()
y_test = test[test.columns[-1]].to_numpy()
print(type(x_test))

```

```
<class 'numpy.ndarray'>
```

```
print(train.columns[-1])
```

```
Class
```


▼ Training and Validation Split

Training and validation set split in 70:30 ratio.

```
x_train, x_val, y_train, y_val = train_test_split(train[train.columns[:-1]].to_numpy(), train
```

```
plt.hist(y_train,color='orange')
```

```
(array([752., 329.,  0., 673.,  0., 283., 345.,  0.,  0., 722.]),
 array([1. , 1.6, 2.2, 2.8, 3.4, 4. , 4.6, 5.2, 5.8, 6.4, 7. ]),
 <a list of 10 Patch objects>)
```



```
print(len(x_train),len(x_val),len(x_test))
```

```
3104 1331 2000
```

▼ TSNE Plot for data visualisaiton

```
from sklearn.manifold import TSNE
```

```
tsne = TSNE(n_components=2,perplexity=50, random_state=6, n_jobs=-1).fit_transform(train)
```

```
tsne_df = pd.DataFrame(tsne)
```

```
tsne_df.sample(10)
```

tsne_df.sample(10)

	0	1
1485	58.583157	-1.797761
258	6.324732	-22.913366
1518	21.538527	45.873734
3715	-15.998539	-30.814167
2325	22.539925	23.036531
2466	34.453903	15.857704
1015	-54.343067	-31.753153
3028	-46.861969	20.278341
1977	40.363178	-7.454031
232	-13.745310	-24.735798

```
plt.figure(figsize=(16,10))
# colorsIdx = {'A': 'rgb(215,48,39)', 'B': 'rgb(215,148,39)'}
# cols      = ['orange','blue']
sns.scatterplot(
    x=0, y=1,
    palette=sns.color_palette("hls", 10),
    data=tsne_df,
    legend="full",
    alpha=1
)
```



<matplotlib.axes._subplots.AxesSubplot at 0x7fe1639d1f90>



```
test_x = np.array([[1,2,3] , [-4,5,-6]])
test_xt = np.array([[0,0,0], [1,2,3], [4,5,6], [-4, 4, -6]])
test_yt = np.array([[1, 1, 2, 3]]).T
```



Double-click (or enter) to edit



▼ kNN algorithm implementation from scratch

```
from scipy.spatial.distance import cdist
```

```
def get_p_y_x_using_knn(y, k, classes):
    first_k_neighbors = y[:, :k]

    N1, N2 = y.shape
    number_of_classes = classes.shape[0]

    prob_matrix = np.zeros(shape=(N1, number_of_classes))

    for i, row in enumerate(first_k_neighbors):
        for j, value in enumerate(classes):
            prob_matrix[i][j] = list(row).count(value) / k

    return prob_matrix
```

```
def predict(X_test, X_train, y_train, k):
    order = cdist(X_test, X_train, metric="euclidean").argsort(kind='mergesort')

    sorted_labels = np.squeeze(y_train[order])
    classes = np.unique(sorted_labels)
    # print(classes)
    p_y_x = get_p_y_x_using_knn(sorted_labels, k, classes)
    number_of_classes = p_y_x.shape[1]
    reversed_rows = np.fliplr(p_y_x)
```

```
prediction = classes[number_of_classes - (np.argmax(reversed_rows, axis=1) + 1)]
return prediction
```

```
from sklearn.metrics import classification_report, accuracy_score
y_predict = predict(x_val, x_train, y_train, 10)
```

```
print(classification_report(y_val, y_predict))
```

	precision	recall	f1-score	support
1.0	0.98	0.97	0.97	320
2.0	0.96	0.94	0.95	150
3.0	0.88	0.93	0.91	288
4.0	0.68	0.59	0.63	132
5.0	0.88	0.82	0.85	125
7.0	0.87	0.90	0.88	316
accuracy			0.89	1331
macro avg	0.87	0.86	0.87	1331
weighted avg	0.89	0.89	0.89	1331

▼ Finding Optimal value of k using grid search

```
# return accuracy for a particular k value when prediction is true by calulating the mean val
```

```
def get_acc(prediction, y_true):
    N1 = prediction.shape[0]
    accuracy = np.sum(prediction == y_true) / N1
    return accuracy
```

```
# returns optimal value of k from a given range of values
```

```
def kselector(x_val, y_val, x_train, y_train, k_values):
    accuracies = []

    for k in tqdm(k_values):
        prediction = predict(x_val, x_train, y_train, k)

        acc = get_acc(prediction, y_val)
        accuracies.append(acc)

    k_optimal = k_values[accuracies.index(max(accuracies))]


    return k_optimal, accuracies
```

```
k_range = np.arange(3,33,2)
best_k, accuracies = kselector(x_val, y_val, x_train, y_train, k_range)
```

100%|██████████| 15/15 [00:08<00:00, 1.84it/s]

```
px.line(y=accuracies, x= k_range, title="Accuracy for k nearest neighbors",  
        labels=dict(x="K Param", y="Accuracy"))
```

Accuracy for k nearest neighbors



```
print("Optimal value of k is:",best_k)
```

```
Optimal value of k is: 5
```

```
y_pred = predict(x_test, x_train, y_train, best_k)
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1.0	0.98	0.99	0.99	461
2.0	0.97	0.96	0.97	224
3.0	0.92	0.92	0.92	397
4.0	0.72	0.69	0.70	211
5.0	0.90	0.87	0.89	237

	7.0	0.85	0.88	0.87	470
accuracy				0.90	2000
macro avg		0.89	0.89	0.89	2000
weighted avg		0.90	0.90	0.90	2000

▼ Comparing our model with sklearn

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
neigh = KNeighborsClassifier(5)
neigh.fit(x_train, y_train)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                     weights='uniform')
```

```
y_pred_skl = neigh.predict(x_test)
```

```
print(classification_report(y_test, y_pred_skl))
```

	precision	recall	f1-score	support
1.0	0.98	0.99	0.99	461
2.0	0.97	0.96	0.96	224
3.0	0.89	0.92	0.91	397
4.0	0.70	0.65	0.67	211
5.0	0.91	0.87	0.89	237
7.0	0.85	0.86	0.86	470
accuracy			0.89	2000
macro avg	0.88	0.88	0.88	2000
weighted avg	0.89	0.89	0.89	2000

```
print("Accuracy for the model from sklearn is",accuracy_score(y_test, y_pred_skl))
print("Accuracy for the model from scratch is",accuracy_score(y_test, y_pred))
```

```
Accuracy for the model from sklearn is 0.894
Accuracy for the model from scratch is 0.9015
```

✓ 1m 19s completed at 2:29 PM ● ✕